

NOSQL

Aggregate data models

- A **data model** is the model through which we perceive and manipulate our data.
- For people using a database, the data model describe how we interact with the data in the database.
- **Data model : the model by which the database organizes data.**

- The **dominant data model** of the last couple of decades is **the relational data model**, which is best visualized as a set of tables.
- Each table has rows, with each **row** representing some **entity** of interest.
- We describe this entity through columns, each having a single value(**atomic value**).
- A column may refer to another row in the different table, which constitutes a relationship between those entities. (**Foreign Key**)

Aggregates

- The relational model takes the information that we want to store and divides it into tuples (rows).
- A tuple is a limited data structure: It captures a set of values, so we cannot nest one tuple within another to get nested records, nor can we put a list of values or tuples within another.
- **Aggregate orientation** takes a different approach.
- We often want to operate on data in units that have a **more complex structure than a set of tuples**.
- **key-value, document, and column-family databases all make use of this more complex record**

- However, there is no common term for this complex record; here (we) use the term “**aggregate.**”
- **aggregate is a collection of related objects that we wish to treat as a unit.**
- In particular, it is a unit for data manipulation and management of consistency.
- Aggregates are also often easier for application programmers to work with, since they often manipulate data through aggregate structures.

NoSQL

NoSQL is a non-relational database management systems, different from traditional relational database management systems in some significant ways.

It is designed for **distributed data stores** where very **large scale of data storing needs**.

For example Google or Facebook which collects terabits of data every day for their users.

These type of data storing **may not require fixed schema**, avoid join operations and typically **scale horizontally**.

RDBMS

- Structured and organized data
- Structured query language (SQL)
- Data and its relationships are stored in separate tables.
- Data Manipulation Language, Data Definition

Language

- Tight Consistency

NoSQL

- Stands for Not Only SQL
- No predefined schema
- Key-Value pair storage, Column Store, Document Store, Graph databases
- Eventual consistency rather ACID property
- Unstructured and unpredictable data
- All NoSQL offerings relax one or more of the ACID properties (will talk about the CAP theorem)

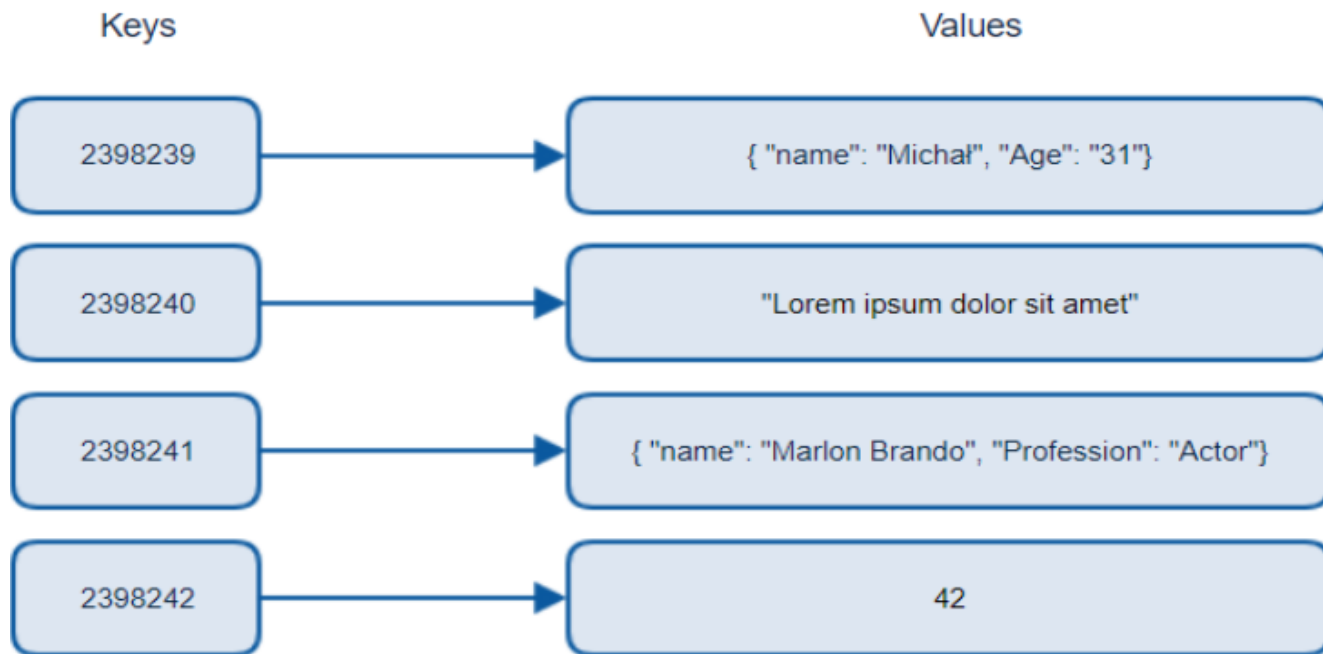
NoSQL Categories

- There are four general types (most common categories) of NoSQL databases.
- Each of these categories has its own specific attributes and limitations.
- There is not a single solutions which is better than all the others, however there are some databases that are better to solve specific problems.
 - Key-value stores
 - Column-oriented
 - Document oriented
 - Graph database

Key-value stores

- Key-value stores are most basic types of NoSQL databases.
- Designed to handle huge amounts of data (Based on Amazon's Dynamo paper).
- Key value stores allow developer to **store schema-less data**.
- In the key-value storage, database stores data as hash table where each key is unique and the value can be string, JSON etc.
- For example a key-value pair might consist of a key like "Name" that is associated with a value like "Robin".
- **Use Cases :**
 - Key-Values stores would work well for shopping cart contents.
 - Customized ad delivery to users based on their data profile
- Example of Key-value store DataBase : **Redis, DynamoDB, Riak**. etc.

- The application is developed on queries that are based on keys.



Implements a **hash table** to store unique keys along with the pointers to the corresponding data values.

Have **no query language** but they do provide a way to **add** and **remove** key-value pairs. Search only based on Key

Column-oriented databases

- Most databases have a row as a unit of storage which, in particular, helps write performance.
- However, there are many scenarios where writes are rare, but we often need to read a few columns of many rows at once for some analytics on those few columns.
- In this situation, it's **better to store groups of columns** for all rows as the basic storage unit—which is why these databases are called column stores.
- Example of Column-oriented databases : BigTable, Hbase, Cassandra etc.

To get a particular customer's name from Figure 2.5 we could do something like `get('1234', 'name')`.

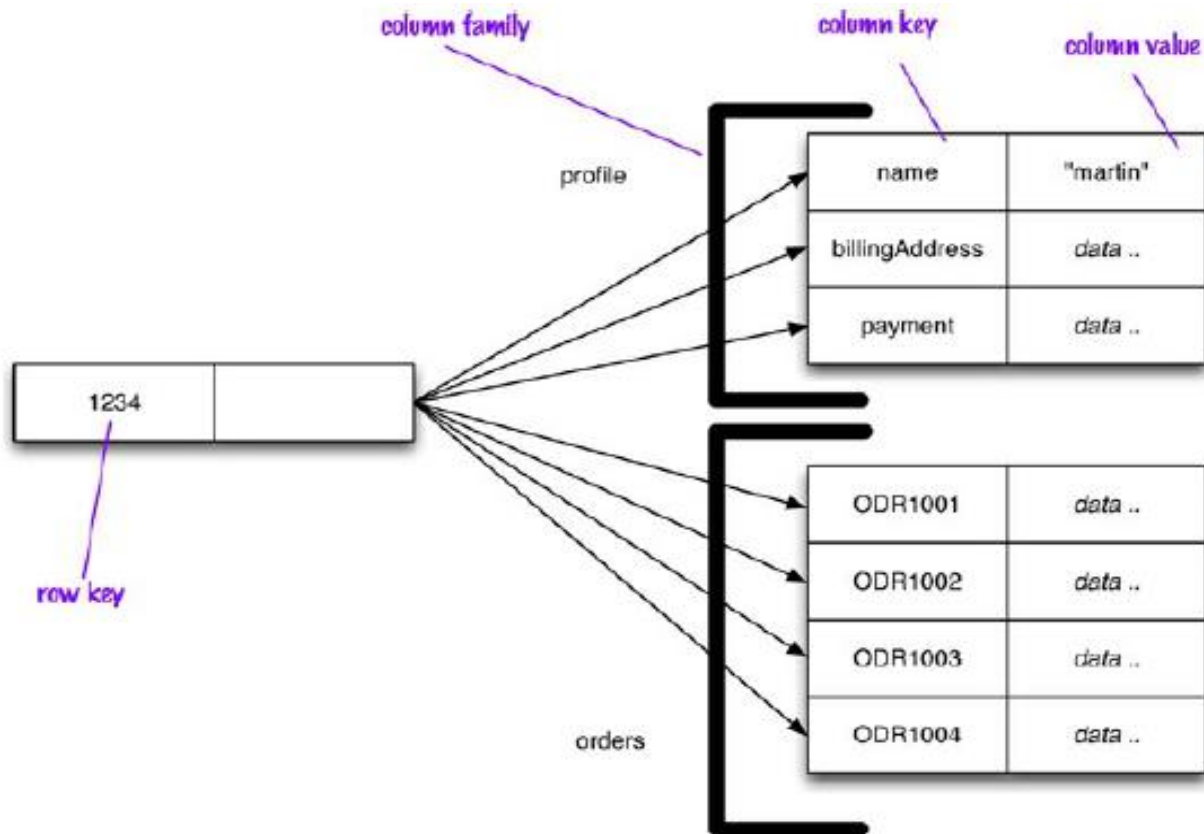


Figure 2.5. Representing customer information in a column-family structure

Document Oriented databases

- A collection of documents
- **Data** in this model is **stored inside documents**.
- A document is a key value collection where the key allows access to its value.
- Documents are **not typically forced to have a schema** and therefore are flexible and easy to change.
- Documents can contain many different key-value pairs, or key-array pairs, or even **nested documents**.
- Example of Document Oriented databases : **MongoDB**, CouchDB etc.

Sammy's contact card document

```
{
  "_id": "sammyshark",
  "firstName": "Sammy",
  "lastName": "Shark",
  "email": "sammy.shark@digitalocean.com",
  "department": "Finance"
}
```

Tom's contact card document with social media accounts information attached

```
{
  "_id": "tomjohnson",
  "firstName": "Tom",
  "middleName": "William",
  "lastName": "Johnson",
  "email": "tom.johnson@digitalocean.com",
  "department": ["Finance", "Accounting"],
  "socialMediaAccounts": [
    {
      "type": "facebook",
      "username": "tom_william_johnson_23"
    },
    {
      "type": "twitter",
      "username": "@tomwilliamjohnson23"
    }
  ]
}
```

Collection

```
{
  "_id": "tomjohnson",
  "firstName": "Tom",
  "middleName": "William",
  "lastName": "Johnson",
  "email": "tom.johnson@digitalocean.com",
  "department": ["Finance", "Accounting"],
  "socialMediaAccounts": [
    {
      "type": "facebook",
      "username": "tomjohnson"
    },
    {
      "type": "twitter",
      "username": "@tomjohnson"
    }
  ]
}
```

```
{
  "_id": "sammyshark",
  "firstName": "Sammy",
  "lastName": "Shark",
  "email": "sammy.shark@digitalocean.com",
  "department": "Finance"
}
```

```
{
  "_id": "tomjohnson",
  "firstName": "Tom",
  "middleName": "William",
  "lastName": "Johnson",
  "email": "tom.johnson@digitalocean.com",
  "department": ["Finance", "Accounting"]
}
```


Graph databases

- A graph database stores data in a graph.
- A graph database is a **collection of nodes and edges**
- Each node represents an entity (such as a student or business) and each **edge** represents a connection or **relationship between two nodes**.
- Every node and edge are defined by a **unique identifier**.
- Each node knows its adjacent nodes.
- Example of Graph databases : OrientDB, **Neo4J**, Titan. etc.

- Graph databases are an odd fish in the NoSQL pond.
- Graph databases are motivated by a different frustration with relational databases and thus have an **opposite model—small records with complex interconnections**, something like Figure below.

- We refer to a **graph data structure of nodes connected by edges**.
- In Figure we have a web of information whose nodes are very small (nothing more than a name) but there is a rich structure of interconnections between them.
- With this structure, we can ask questions such as “ **find the books in the Databases category that are written by someone whom a friend of mine likes.**”