

Pipes

- **Pipes represent a channel for InterProcess Communication.**
- **Two types are: Unnamed and Named Pipes(FIFO)**

Pipe

- A simple, **unnamed** pipe provides a one-way flow of data.
- An unnamed pipe is created by calling *pipe()*, which returns an array of 2 file descriptors (int).
 - The file descriptors are for reading and writing, respectively

pipe System Call (unnamed)

Creates a half-duplex pipe.

- Include(s): `<unistd.h>`
- Syntax: `int pipe (int pipefd[2]);`
- Return: Success: 0; Failure: -1; Sets errno: Yes
 - What does it mean to return errno?
- If successful, the *pipe* system call will return two integer file descriptors, `pipefd[0]` and `pipefd[1]`.
 - `pipefd[1]` is the write end to the pipe.
 - `pipefd[0]` is the read end from the pipe.
- Parent/child processes communicating via unnamed pipe.

Features of Pipes

Features of Pipes


- On many systems, pipes are limited to 10 logical blocks, each block has 512 bytes.
- As a general rule, one process will write to the pipe (as if it were a file), while another process will read from the pipe.
- Data is written to one end of the pipe and read from the other end.
- A pipe exists until both file descriptors are closed in all processes

Piping Between Two Processes

- The pipe is represented in an array of 2 file descriptors (int)

Writing process

Reading process




Full Duplex Communication via Two Pipes

Two separate pipes, say $p0$ and $p1$

Process A

Process B



write System Call

Function:

- To write *nbytes* to the write end of a pipe.
 - If a write process attempts to *write* to a full pipe, the default action is for the system to block the process until the data is able to be received.
- Include(s): `<unistd.h>`
- Syntax: *ssize_t write (int fd, const void *buf, size_t nbytes);*
 - just the *write* system call
- Returns
 - success: Number of bytes written; Failure; -1; Sets errno: Yes.
- Arguments
 - *int fd*: file descriptor;
 - *const void *buf*: buffer;
 - *size_t nbyte*: number of bytes in buffer

read System Call

Function:

- To read *nbytes* from the read end of a pipe.
 - if *read* is attempted on an empty pipe, the process will block until data is available.
- Includes: `<unistd.h>` `<sys/types.h>` `<sys/uio.h>`
- Syntax: *ssize_t read(int fd, const void *buf, size_t nbytes);*
- Return
 - success: Number of bytes read;
 - Failure; -1; Sets errno:Yes.
 - EOF (0): write end of pipe closed
- Arguments
 - *int fildes*: file descriptor;
 - *const void *buf*: buffer;
 - *size_t nbyte*: number of bytes

Unnamed Pipes

- Unnamed pipes can only be used between related process, such as parent/child, or child/child process.
- Unnamed pipes can exist only as long as the processes using them.
- An unnamed pipe is constructed with the *pipe* system call.

Named pipes

Named pipes are used for inter-process communications.

Features:

- Exist as special files in the physical file system
- Any unrelated processes can access a named pipe and share data through it
- Access to named pipes is regulated by the usual file permissions
- Pipe data is accessed in a FIFO style
- Once created, a named pipe remains in the file system until explicitly deleted

Creation:

- By UNIX shell commands. Example:

```
mknod <filename> p  
mkfifo a=rw <filename>
```

- By systems calls. Example:

```
mknod(char *pathname, mode_t mode, dev_t dev);  
mknod("/tmp/myfifo", S_IFIFO | 0660, 0);
```

Named pipes (cont.)

Same I/O operations style on named pipes and regular files – *open()*, *read()* and *write()* calls.

Implementation of I/O operations:

- By system calls.
- By library functions.

Semantics of *open()* call:

- **Blocking.** The process that opens the named pipe for reading, sleeps until another process opens it for writing, and v.v.
- **Non-blocking.** Flag *O_NONBLOCK*, used in *open()* call disables default blocking.


Pipes have size limitations.

Named pipes. Example

Client-server communication through a pipe

Server

```
#include <fcntl.h>
...
#define PIPE "fifo"
int main(){
    int fd;
    char readbuf[20];
    mknod(PIPE, S_IFIFO | 0660, 0); // create pipe
    fd = open(PIPE, O_RDONLY, 0);    // open pipe
    for (;;) {
        if (read(fd, &readbuf, sizeof(readbuf)) < 0){ //read from pipe
            perror("Error reading pipe");
            exit(1);
        }
        printf("Received string: %s\n", readbuf); // process data
    }
    exit(0);
}
```



Named pipes. Example (cont.)

Client

```
#include <stdio.h>
...
#define PIPE "fifo"

int main(){
    int fd;
    char writebuf[20] = "Hello"; // open pipe
    fd = open(PIPE, O_WRONLY, 0);
    // write to pipe
    write(fd, writebuf, sizeof(writebuf));
    exit(0);
}
```

Redirecting Standard I/O

- A process that communicates solely with another process doesn't use its standard I/O.
- If process communicates with another process only via pipes, redirect standard I/O to the pipe ends
 - Functions: *dup*, *dup2*

dup & dup2

```
#include <unistd.h>
int dup(int files);
```

- Returns a new file descriptor that is a copy of *files*
- File descriptor returned is first available file descriptor in file table.
- For example, to dup a read pipe end to stdin (0), close stdin, then immediately dup the pipe's read end.
- Close unused file descriptors; a process should have only one file descriptor open on a pipe end.

dup & dup2

```
#include <unistd.h>
```

```
int dup2(int fromFD, int toFD);
```

- Duplicate *fromFD* to *toFD*. If *toFD* is open, it is closed first.
- For example, if a pipe's ends are in array *pipefd*, ***dup2(pipefd[1],1)*** redirects stdout to the write end of the pipe.
- You still must close the unused pipe end, in this case *pipefd[1]*